

Spark DataFrame:

- **DataFrame** is a **distributed collection** of data organized into **named columns**. It is **conceptually** equivalent to a **table** in a **relational database**.
- DataFrame **appeared** in **Spark Release 1.3.0**.
- The idea behind DataFrame is it **allows processing** of a **large amount** of **structured data**. DataFrame contains rows with **Schema**.
- **DataFrame** in **Spark** **overcomes RDD** but **contains** the **features** of **RDD** as well.
- The features **common** to **RDD** and **DataFrame** are **immutability**, **in-memory** and **distributed computing capability**.
- It **allows** the **user** to **impose** the **structure** onto a **distributed collection** of data.
- We can **build DataFrame** from **different data sources**. For e.g. **structured data file**, **tables in Hive**, **external databases** or **existing RDDs**.

DF Reference to structure data like RDBMS:

Timestamp	Age	Gender	Country	state
2014-08-27 11:29:31	37	Female	United States	IL
2014-08-27 11:29:37	44	M	United States	IN
2014-08-27 11:29:44	32	Male	Canada	
2014-08-27 11:29:46	31	Male	United Kingdom	
2014-08-27 11:30:22	31	Male	United States	TX

Data Frame Schema

1. Column Names
2. Data Types

Limitations of Spark RDD:

- It **does not** have any **built-in optimization engine**.
- There is **no provision** to **handle structured data**.
- It **does not** have **schema**.
- Thus to **overcome** these **limitations**, **DataFrame** came into existence.

Creating DataFrame in Apache Spark:

To **all** the **functionality** of **Spark**, **SparkSession** class is the **entry point**. For the **creation** of basic **SparkSession** just use:

sc → **SparkContext**
spark → **SparkSession**

Using **SparkSession**, an **application** can **create DataFrame** from an **existing RDD**, **Hive table** or from **Spark data sources**. Using **Spark SQL DataFrame** we can **create a temporary view** and run **SQL queries** on the **data**.

Convert RDD to DF:

- In PySpark, **toDF()** function of the **RDD** is used to convert **RDD** to **DataFrame**.
- We **need** to **convert RDD** to **DF** as **DF** provides **more advantages** over **RDD**.
- **DF** is a **distributed collection** of **data organized** into **named columns** similar to **DB tables** and provides **optimization** and **performance improvements**.

First, let's **create** an **RDD** by **passing Python list object** to **sparkContext.parallelize function**. We would need this "rdd" object for all our examples below.

A **list** is a **data structure** in **Python** that's **holds** a **collection** of **items**. **List items** are **enclosed** in **square brackets**, like this *[data1, data2, data3]*.

In PySpark, when you have **data** in a **list meaning** you have a **collection** of **data** in a **PySpark driver memory** when you **create** an **RDD**, this **collection** is **going to** be **parallelized**.

1) Using rdd.toDF () function:

PySpark **provides toDF ()** function in **RDD** which can be used to convert **RDD** into **Dataframe**.

Note: By default, **toDF ()** function **creates column names** as **"_1"** and **"_2"**. This snippet yields below schema.

toDF () has **another signature** that **takes arguments** to **define column names** as shown below.

2) Using PySpark createDataFrame () function:

SparkSession class provides **createDataFrame ()** method to **create DataFrame** and it takes **rdd object** as an **argument** and **chain** it with **toDF ()** to specify **names** to the **columns**.

3) Using createDataFrame () with StructType schema:

- When you **infer** the **schema**, by **default** the **datatype** of the **columns** is **derived** from the **data** and set's **nullable** to **true** for **all columns**.
- We can **change** this **behavior** by **supplying schema** using **StructType** – where we can **specify** a **column name**, **data type** and **nullable** for each **field/column**.

4) Create DataFrame from Data sources:

- In **real-time** mostly you **create DataFrame** from **data source files** like **CSV**, **Text**, **JSON**, and **XML** etc.
- PySpark **provides csv ("path")** on **DataFrameReader** to **read** a **CSV file** into **PySpark**

DataFrame and **dataframeObj.write.csv** ("path") to **save** or **write** to the **CSV** file.

- **PySpark** **supports** reading a **CSV** file with a **pipe**, **comma**, **tab**, **space**, or any **other delimiter/separator** files.

PySpark Read CSV file into DataFrame:

- Using **csv** ("path") or **format** ("csv").**load** ("path") of **DataFrameReader**, you can **read** a **CSV** file into a **PySpark DataFrame**.
- These **methods** take a **file path** to **read from** as an **argument**.

a) Reading CSV file & see the difference in data for headers & schema

Note: This example **reads** the **data** into **DataFrame** columns **"_c0"** for the **first column** and **"_c1"** for the **second** and **so on** and by **default data type** for **all these columns** is **treated** as **String**.

Using Header record for column names:

This **option** is **used** to **read** the **first line** of the **CSV** file as **column names**.

By **default** the **value** of this **option** is **False**, and **all column types** are **assumed** to be a **string**. **Not mentioning** this, the **API** **treats** **header** as a **data record**.

b) Reading the Header & see the difference in data:

c) delimiter:

delimiter **option** is **used** to **specify** the **column delimiter** of the **CSV** file.

By **default**, it is **comma (,)** character, but can be **set** to any character like **pipe(|)**, **tab (\t)**, **space** using this **option**.

d) inferSchema:

The **default value** **set** to this **option** is **False** when **setting** to **true** it **automatically** **infers column types** based on the **data**. Note that, it **requires** **reading** the **data** **one more** time to **infer** the **schema**.

e) Read multiple CSV files:

Using the **read.csv()** **method** you can also **read multiple csv** files, just pass **all file names** by **separating comma** as a **path**, for example:

f) Read all CSV files in a directory:

We can **read all CSV** files from a **directory** into **DataFrame** just by **passing directory** as a **path** to the **csv ()** **method**.

g) Reading CSV files with a user-specified custom schema:

If you **know** the **schema** of the **file** **ahead** and **do not** want to **use**

the **inferSchema** option for **column names** and **types**, use **user-defined custom** column names and **type** can be given using **schema** option.

h) Write PySpark DataFrame to CSV file:

Use the **write()** method of the PySpark **DataFrameWriter** object to **write** PySpark **DataFrame** to a **CSV** file.

While **writing** a **CSV** file you can **use several options**.

E.g. header to **output** the **DataFrame** **column names** as **header record** and **delimiter** to **specify** the **delimiter** on the **CSV** **output** **file**.

Saving modes:

PySpark DataFrameWriter also has a **method mode ()** to **specify saving mode**.

- 1) error:** This is a **default option** when the **file already exists**, it **returns an error**.
- 2) ignore:** **Ignores write operation** when the **file already exists**.
- 3) append:** To **add the data** to the **existing file**.
- 4) overwrite:** This **mode is used** to **overwrite** the **existing file**.

i) Select single & multiple columns from PySpark:

- You can **select** the **single** or **multiple column** of the **DataFrame** by **passing** the **column names** you **wanted** to **select** to the **select ()** **function**.
- Since **DataFrame's** are **immutable**, this **creates** a **new DataFrame** with a **selected columns**. **show ()** **function** is **used** to **show** the **Dataframe** **contents**.

Single & Multiple Columns:

```
df.select("firstname").show()  
df.select("firstname", "lastname").show()
```

Using Dataframe object name:

```
df.select(df.firstname, df.lastname).show()
```

Using col function:

```
df.select(col("firstname"), col("lastname")).show()
```

j) Select nested struct columns from PySpark:

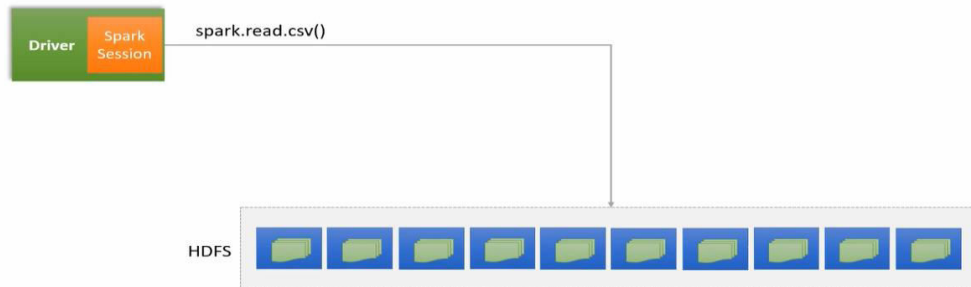
If you have **struct (StructType)** **column** on **PySpark DataFrame**, you **need** to **use** an **explicit column** **qualifier** in **order** to **select**.

k) In order to **get** the **specific column** from a **struct**, you **need** to **explicitly** **qualify**.

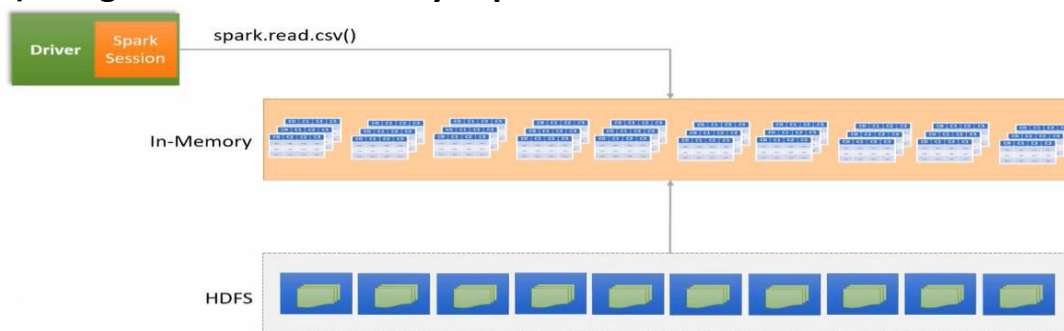
l) In order to **get** all **columns** from **struct** **column**.

DataFrame Partitions & Executors:

1) Distributed Storage in multiple nodes:

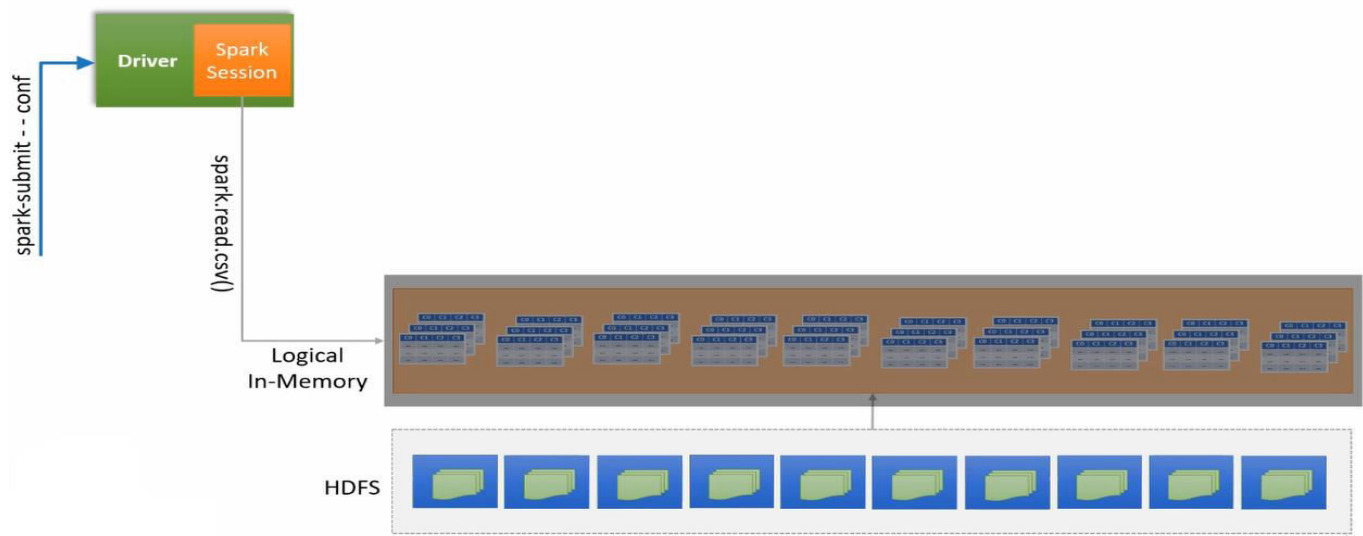


2) Brings the data in-memory in partitions:



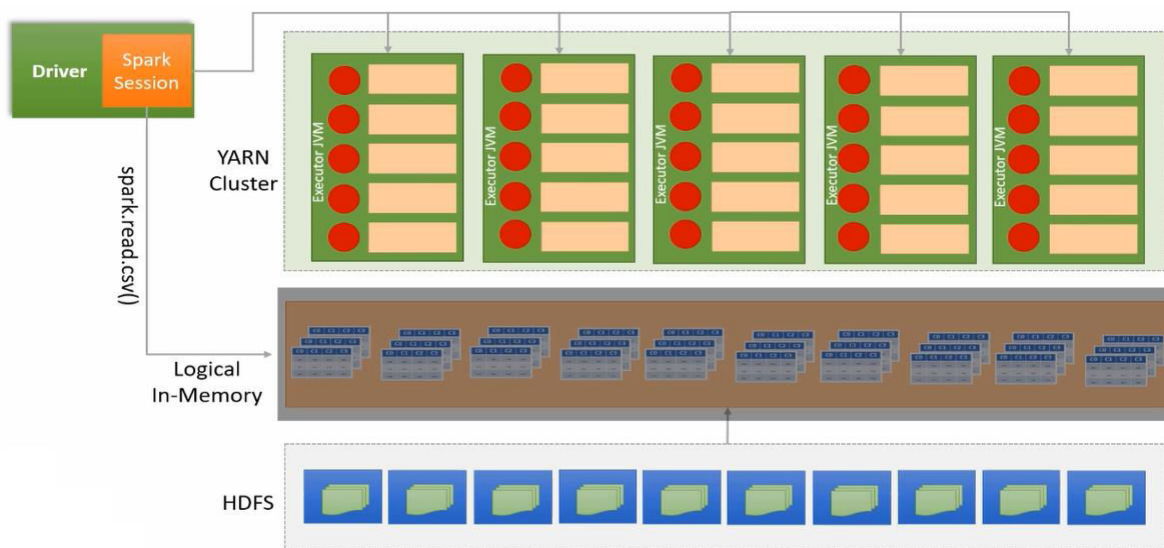
3)

- Driver reaches out to SM & CM to get the details of file partitions.
- So at runtime your driver know how to read the data files & how many partitions are there.
- So it creates logical in-memory structure which we see as a DF.



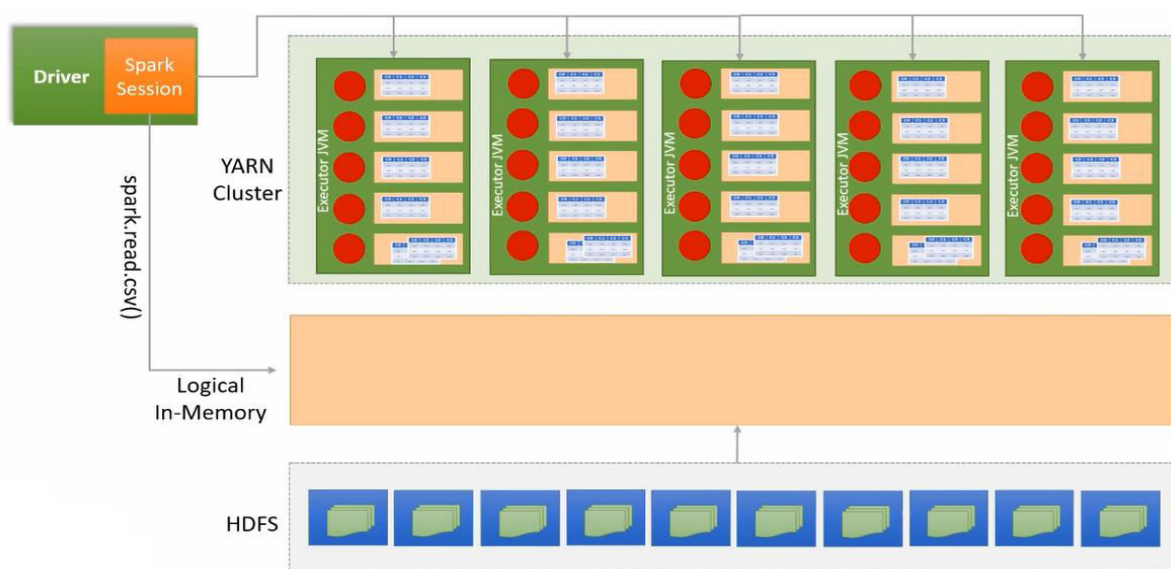
4)

- We can calculate how many executors, memory can be allocated to these executors which is given in your spark-submit.
- So all these configurations are available to your driver. Let's assume we configured to start 5 Executors each with 10GB memory and 5 Cores.
- Now the driver again will reach out to CM and ask for containers. One those containers are allocated the driver starts executors within these containers.
- Each executor is nothing but a JVM process with some assigned cores & memory.
- So here each executor is started with 5 executor cores & 10GB memory.



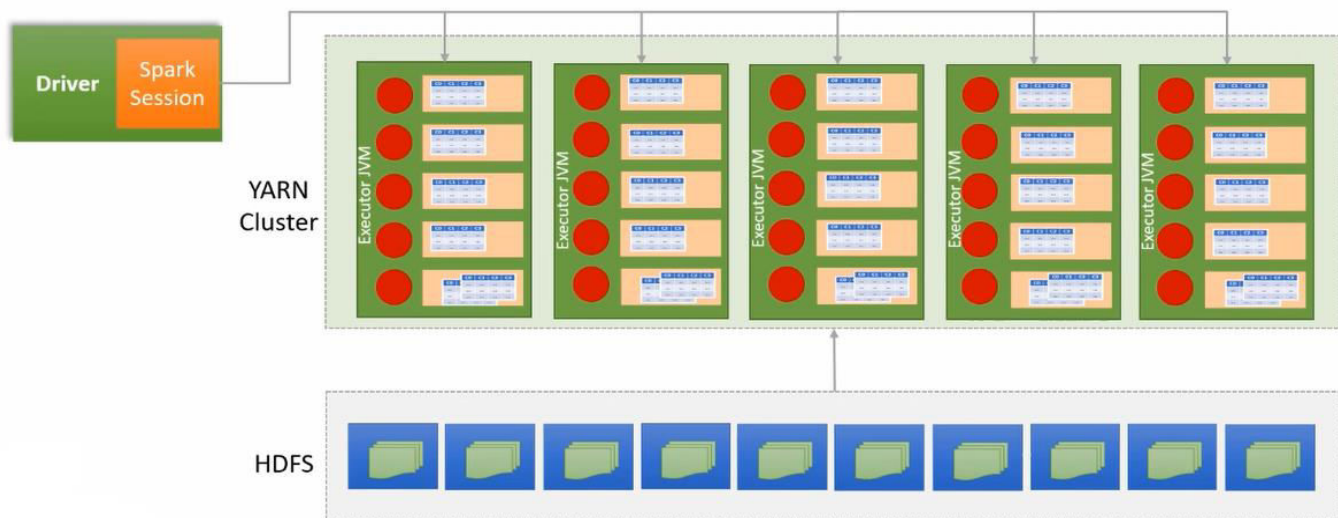
5)

- Now the Driver is ready to distribute the work to these executors.
- So driver assigns some DF partitions to each JVM core.
- These executor core will load their respective partitions in-memory.



6)

- Now you are ready with your distributed dataframe setup where each executor core is assigned its own data partition to work on.
- In all this process spark will also try to minimize the network bandwidth for loading data from physical storage to the JVM memory. How? That's the internal spark optimization.
- While assigning partition to these executors spark will try its best to allocate the partitions which are closed to the executors to the network.
- However such data locality is not always possible so spark & CM will work together to achieve best possible localization.




Transformations and Actions:

a) Transformations:

- Transformations are operations which will transform your RDD data from one form to another.
- When you apply this operation on any RDD, you will get a new RDD of transformed data (RDDs in Spark are immutable).
- Operations like **map**, **filter** and **flatMap** are transformations.

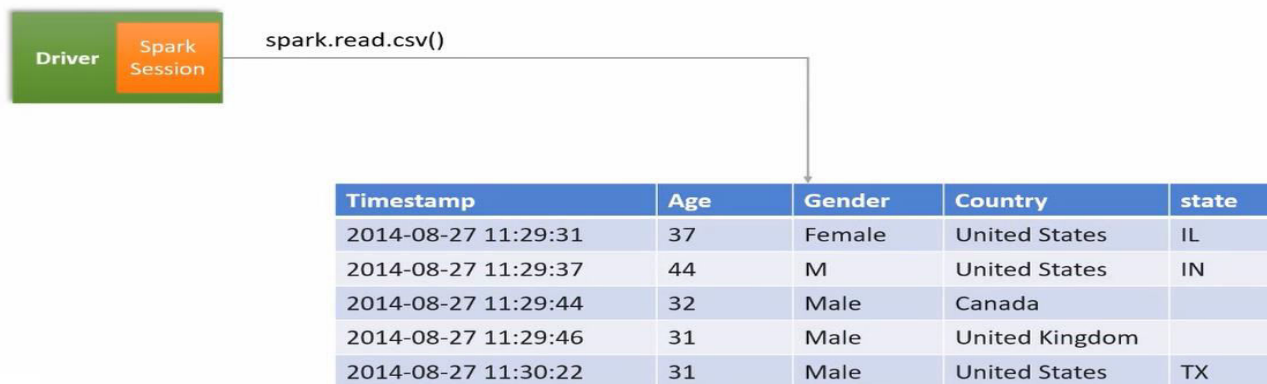
b) Actions:

- When action is triggered new RDD is not formed like transformations. Thus, actions are operation that gives non-RDD values.
- The values of action are stored to drivers or to the external storage system.
- It brings laziness of RDD into motion.
- An action is one of the ways of sending data from executor to the driver.
- Operations like **show**, **read**, **write**, **count**, **collect**, **save** are actions.



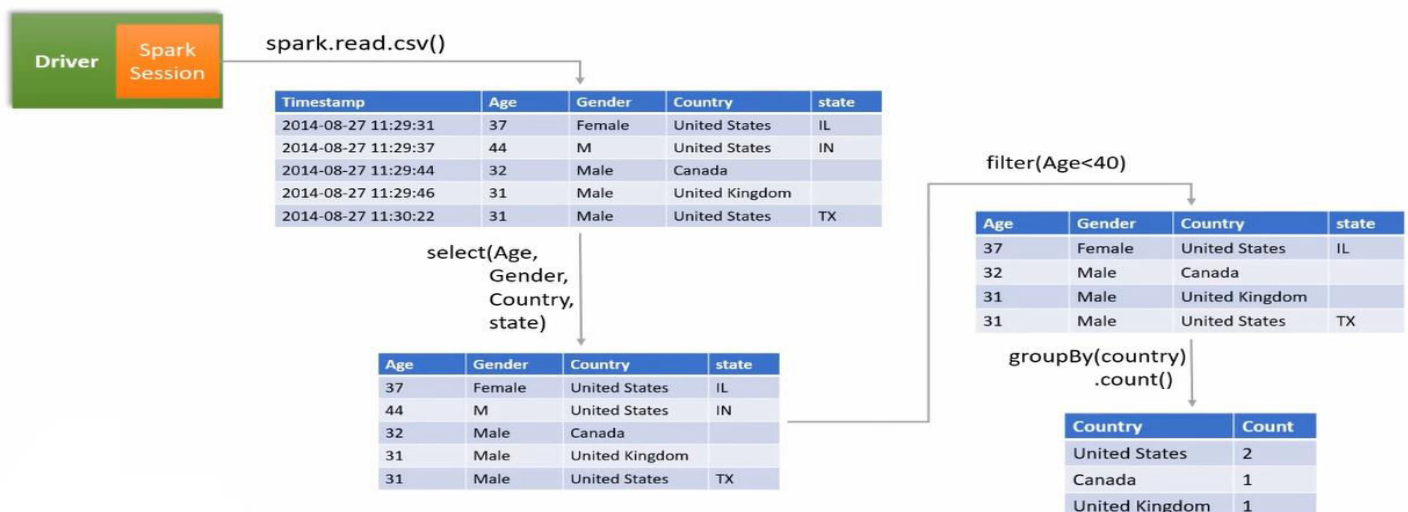
Transformations (<i>lazy</i>)	Actions
select	show
distinct	count
groupBy	collect
sum	save
orderBy	
filter	
limit	

1)
When we read DF it looks like simple DF by hiding all the complex data distribution.



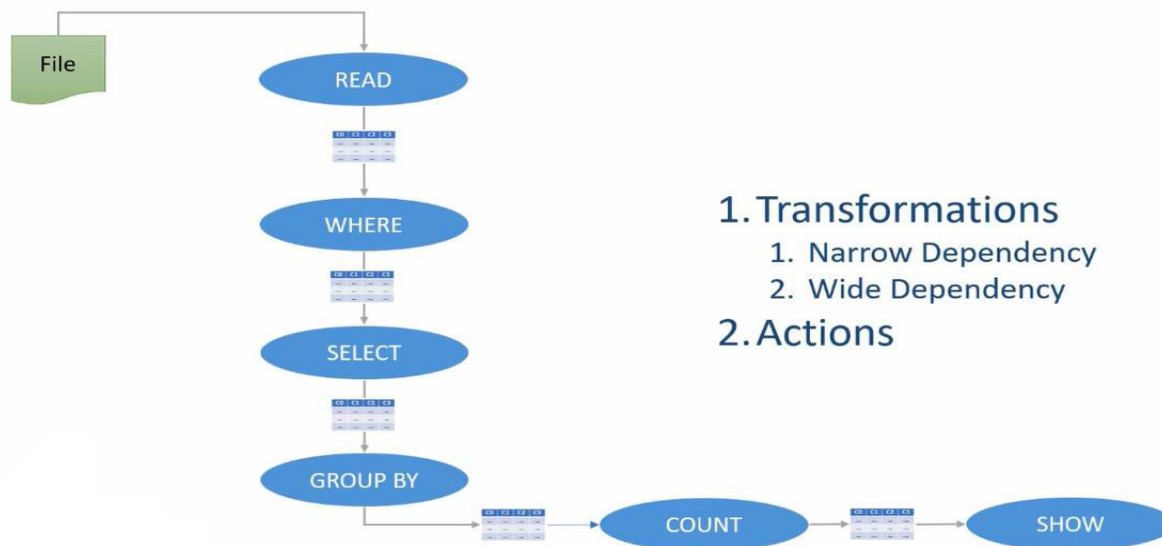
2) Immutability:

- Spark DF are immutable however you can give instructions to your driver what you want to do & let driver decide how to achieve it with the executors.
- These instructions to the driver are called transformations & they can be select, filter, groupBy.



3)

Graph of transformation operations which will create DAG.s



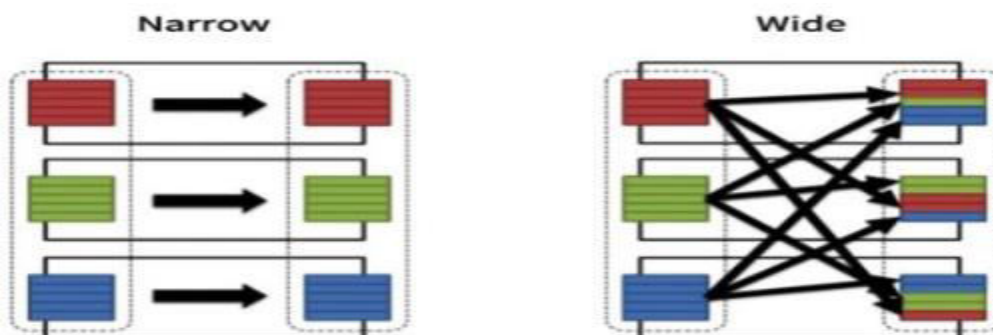
There are two types of transformations: **Narrow and Wide.**

a) Narrow: In Narrow transformation all the elements that are required to compute the records in a single partition lives in the single partition of parent RDD.

E.g. map (), filter (), coalesce () etc.

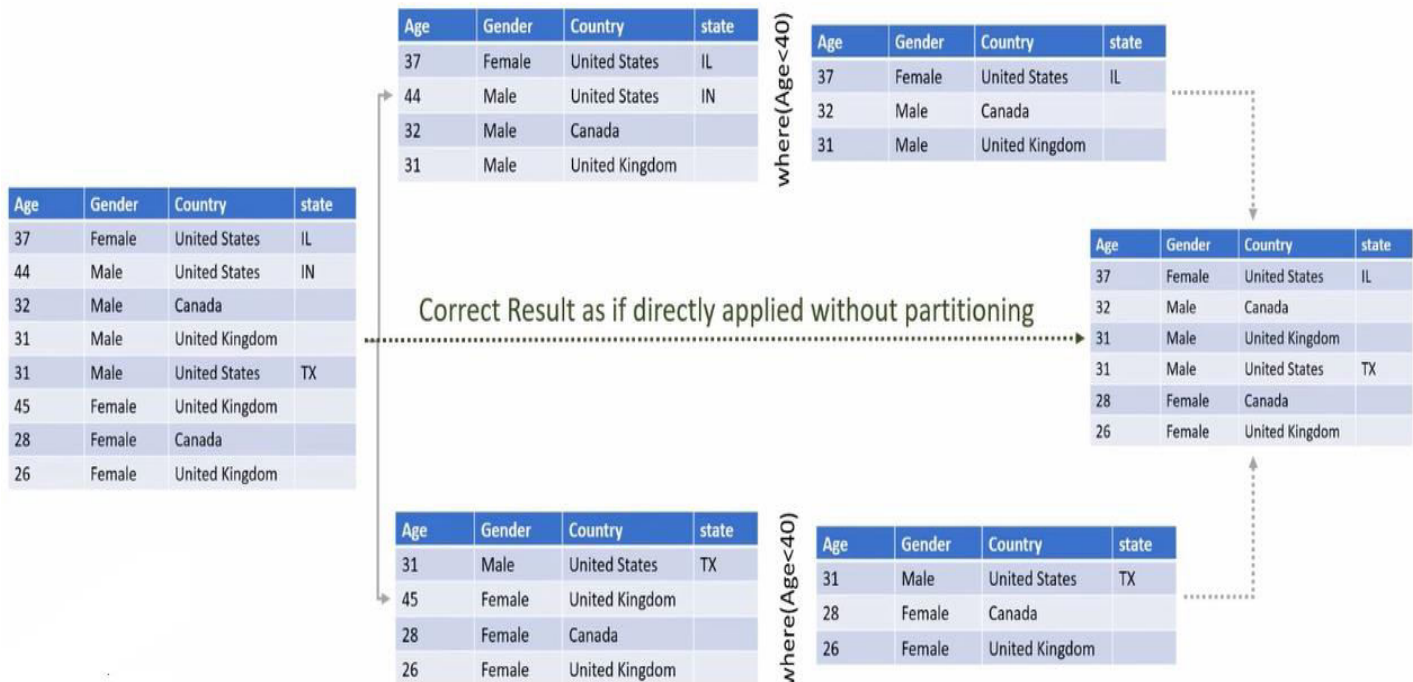
b) Wide: In wide transformation all the elements that are required to compute the records in a single partition lives in many partitions of parent RDD.

E.g. distinct (), groupBy (), repartition () etc.



4) Narrow Transformation:

A transformation performed independently on a single partition to produce valid results.



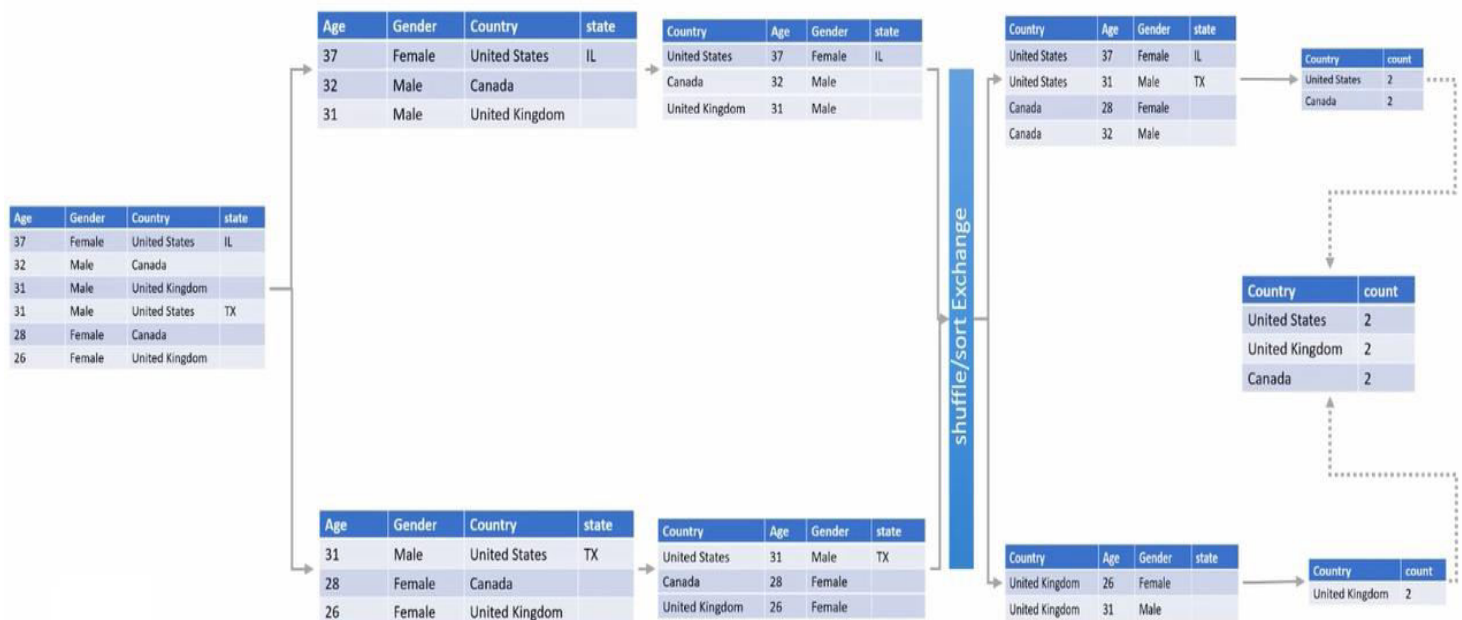
5) Wide:

A transformation that requires data from other partitions to produce valid results.

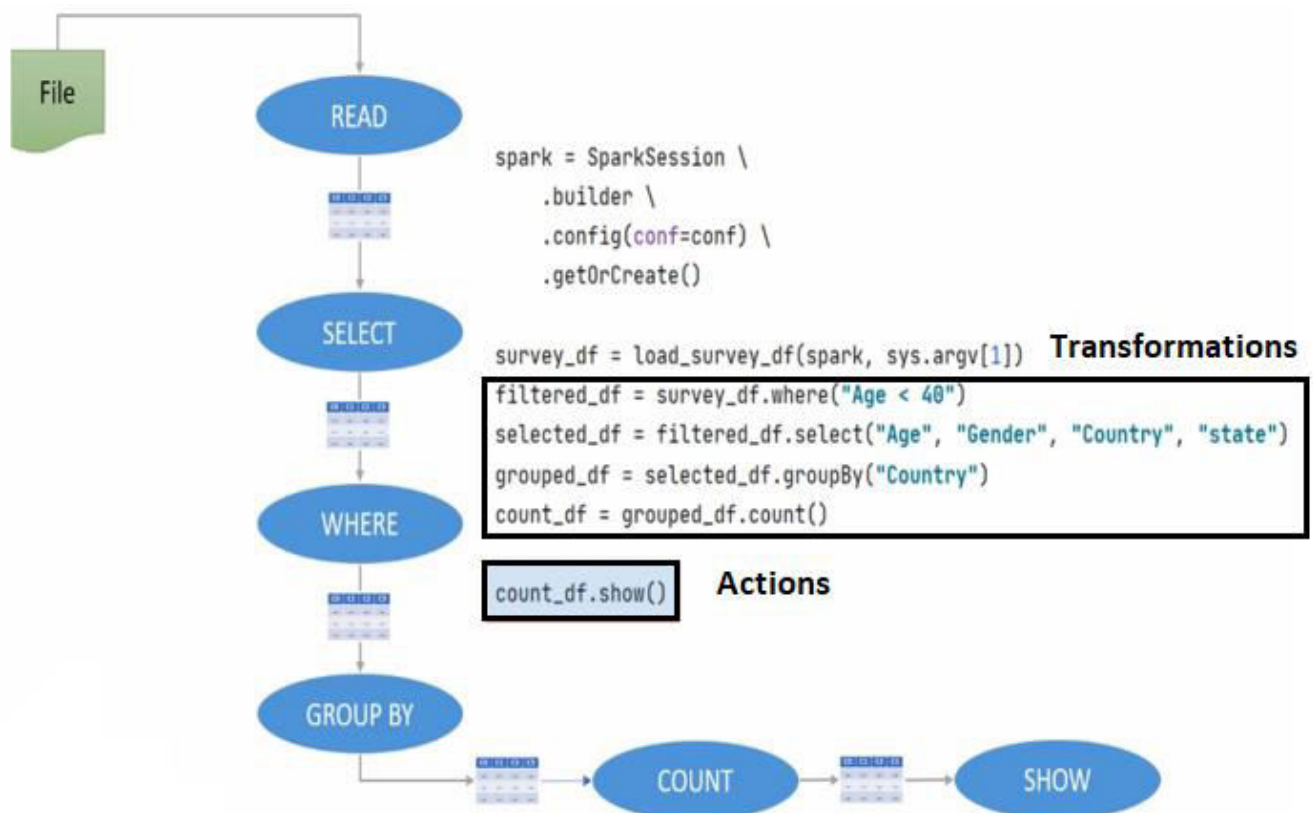


6)

A transformation that requires data from other partitions to produce valid results.



7) Lazy Evaluation

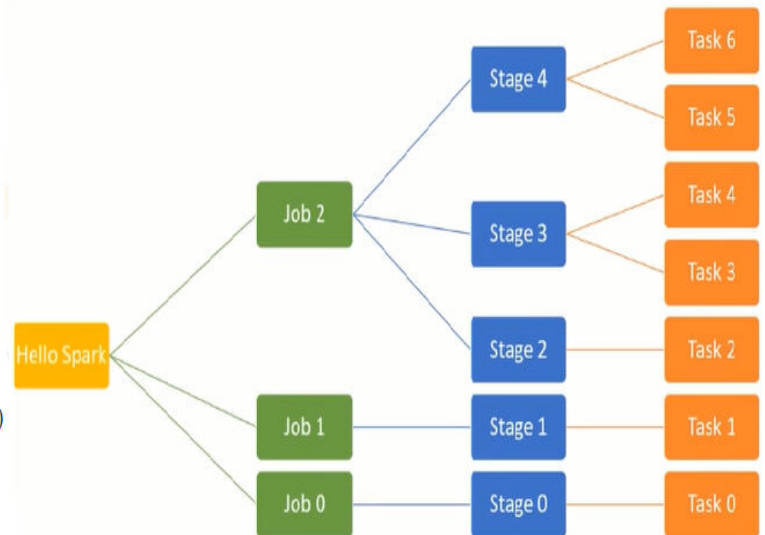


Spark Jobs, Stages & Tasks:

```

4
5 ▶ if __name__ == '__main__':
6     spark = SparkSession.builder \
7         .master("local[3]") \
8         .appName("My First Program") \
9         .config("spark.sql.shuffle.partitions", "2") \
10        .getOrCreate()
11
12    df = spark.read.format("csv") \
13        .option("header", "true") \
14        .option("inferSchema", "true") \
15        .load(input_path)
16    partition_df = df.repartition(2)
17
18    filter_df = partition_df.where("Age < 40")
19    select_df = filter_df.select("Age", "Gender", "Country", "State")
20    group_df = select_df.groupBy("Country")
21    count_df = group_df.count()
22
23    count_df.collect()
24    input("Enter to Quit")
25

```

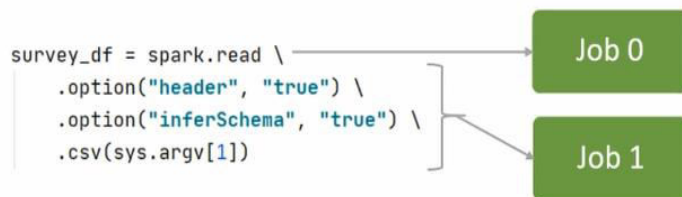


Understanding Your Execution Plan:

Step 1: Only Read csv file → 1 Job → 1 Stage → 1 Task

Step 2: Read Header & inferSchema → 1 Job 1 Stage → 1 Task

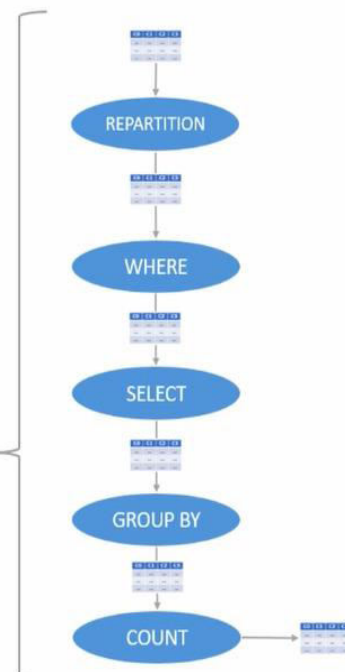
Step 3: Collect DF → 2 Jobs → 3 Stages → Tasks



```

partitioned_survey_df = survey_raw_df.repartition(2)
count_df = survey_df.filter("Age < 40") \
    .select("Age", "Gender", "Country", "state") \
    .groupBy("Country") \
    .count()
logger.info(count_df.collect())

```



DAG for collect action:

